

30d (Modeling and Simulation)
DOE FY18, Release 2 Phase II SBIR/STTR

Phase II Proposal

VnV: A Self Documenting Testing Framework for In-situ
Verification and Validation in High Performance Computing
Applications.



Phase I Contract Number: DE-SC0018728

Ben O'Neill¹, PI

Gregory Watson²

¹ RNET Technologies
240 W Elmwood Dr
Dayton, OH 45459-4296
boneill1@RNET-Tech.com

² Oak Ridge National Laboratory
1 Bethel Valley Road MS6016
Oak Ridge, TN 37831
watsongr@ornl.gov

  “Pages ~~1-2~~ **FO** and **TODO** of this document may contain trade secrets or commercial or financial information that is privileged or confidential and is exempt from public disclosure. Such information shall be used or disclosed only for evaluation purposes or in accordance with a financial assistance or loan agreement between the submitter and the Government. The Government may use or disclose any information that is not appropriately marked or otherwise restricted, regardless of source. Proprietary information is marked with blue curly brackets (i.e., { ... }).”

Contents

1 Significance, Background, and Technical Approach

1.1 Identification and Significance of the Innovation

RNET Technologies Inc. (RNET) in Dayton, OH and Oak Ridge National Laboratory (ORNL) are responding to 2019 DOE SBIR/STTR Phase II Release 2 (DE-FOA-0001976). This proposal is for a Phase II contract in succession to an initial Phase I contract (Contract #: DE-SC0018728) awarded for **topic** DOE SBIR/STTR Topic 30d (Modeling and Simulation). Based on a prototype developed in Phase I, RNET and ORNL are proposing the development of VnV: a self-documenting testing framework for in-situ **solution** verification and validation in high performance computing applications. In this proposal we will highlight

the need for, and the tremendous value of, the proposed tool in high performance numerical simulations **ns like those in the NEAMS toolkit**. The goal of the proposed Phase II project will be to develop and demonstrate a production quality implementation of the VnV framework that can deliver that value to a broad spectrum of numerical simulation applications.

The overarching goal of the VnV toolkit is to facilitate the development of *explainable* numerical simulations that provide the user with not only the final solution, but also an explanation of how the solution was obtained and why it can be trusted. In doing so, the VnV toolkit will streamline the process of solution verification and validation for developers and end-users alike. RNET has extensive SBIR experience in various aspects of High Performance Computing such as performance optimization of numerical softwares and libraries, development of fine-grained power monitoring tools for HPC infrastructure, and the development of software usability tools that enhance the user experience when working with HPC simulations. Examples include a machine learning based plugin for automated linear solver selection in NEAMS tools and a cloud based workflow management tool that allows users to design, execute and visualize simulations on remote machines through a standard web browser. Dr. Watson (ORNL) has extensive experience developing parallel debugging software for high performance computing applications. **Dr. Watson is an experienced software professional with extensive knowledge of software design, architecture, and engineering practices.** He is the project leader of the open-source Eclipse Parallel Tools Platform project and has expertise in the development of highly scalable tools and communication frameworks for peta-scale high performance computing systems. We believe that RNET's experience developing advanced numerical software for the HPC community combined with Dr. Watson's experience with scientific software engineering uniquely qualifies our team to fully develop and commercialize the VnV framework.

1.1.1 Identification and Significance

Numerical modeling and simulation (M&S) is almost always more economical than live testing and prototyping; a fact that has seen wide-scale uptake of numerical **M&Sin** the R&D **lifecycle** of products ranging from \$10 polycarbonate **kids** toys, up to solar panels, airplane wings and nuclear reactors. With access to large-scale computational resources at an all time high, and with exascale computing resources on the horizon; the role of numerical **M&Sin** the design of next generation technologies is only expected to increase. However, numerical simulations are, by definition, an *approximation* to a real world physical system. As such, it is important that this increased reliance on simulated tests is accompanied **by concerted effort** to ensure simulations are fit for the intended purpose. According to the DoD best practices **guide**, verification and validation (V&V) of a code should be performed when *...the risk of making an incorrect decision based on a simulation outweighs the time and cost of performing V&V to ensure that a simulation produces results that are sufficiently accurate and reliable.* One only needs to look to the Sleipner platform accident, where an offshore platform collapsed due to failures in the finite element simulation, to get an idea of the devastating consequences poorly verified numerical simulations can have on business, the environment and society.



In recent years, the DOE and other government agencies has been pushing for increased technology transfer between government funded research institutions (both at national laboratories and academic institutions) and industry. For numerical simulation technologies, technology transfer often manifests as the release of software using a license that contains language to the effect of **use at your own risk, even** in cases where the software has been rigorously verified and validated throughout the development life-cycle. **In other words, it** is ultimately the responsibility of the end user to ensure the inputs **provided to the simulation package lies** within the scope of the applications capabilities, and that the solution obtained is an accurate representation of the physical model of interest to them. After all, the direct costs of a design failure, be it time, money, or loss of life, fall squarely on the shoulders on the products creator, and any attempt to shift the blame to the developers of simulation library *X* will likely fall of deaf ears. **While** it is impractical to unilaterally assert the accuracy of a complex simulation for any given user input, there is a significant need

for tools that aid the end-user in the process of verifying and validating solutions obtained from general purpose simulation packages. Such tools will significantly reduce the overheads associated with incorporating high performance M&Stools in the design pipelines of emerging technologies.

End-user V&Vis of particular importance in the DOE Nuclear Energy Advanced Modeling and Simulation (NEAMS) program. NEAMS is developing predictive models for the advanced nuclear reactor and fuel cycle systems using leading edge computational methods and high performance computing technologies **?**. An important objective of the NEAMS program is to enable widespread use among the industry, academia, and regulatory communities[?]. This objective has lead to the development of the NEAMS workbench, which has significantly increased the overall user experience of the NEAMS tools. The NEAMS group **has placed** a significant emphasis on V&Vin the NEAMS toolkit, as outlined in the NEAMS V&Vplan (version 0). **The result of this is** that the NEAMS tools have integrated functionality for input validation (through the workbench and MOOSE), for performing mesh refinement and method of manufactured solution analysis (through MOOSE) and for completing uncertainty quantification and sensitivity analysis with DAKOTA (through the workbench). Despite this, the complex nature of the codes, combined with the expert knowledge required to set up V&Vtesting and the infinite array of input parameters available in some tools (PETSc, the linear solver package used in MOOSE has thousands of configuration options alone) makes setting up robust end-user V&Van overwhelming task for all but the most expert users.

In order to address the difficulties with end user solution verification in the NEAMS toolkit, and the numerical simulation community as a whole, RNET Technologies Inc. (RNET) and Oak Ridge National Laboratory (ORNL) are proposing the development of the VnV framework; a C/C++ software package that streamlines end user V&V. The framework will focus on providing robust functionality for the V&Vof solutions obtained from general purpose numerical simulation packages (e.g., MOOSE, PETSc, libMesh, Fenics, OpenFoam, etc). To do this, the framework will facilitate the development of *explainable* numerical simulations that, in addition to producing a simulation result, produce a detailed report outlining how the solution was obtained and why it should be trusted. From the perspective of the DOE, the explainable numerical simulations promoted by the proposed framework will increase the appeal its government funded simulation technologies, drive technology transfer and help ensure the nation benefits from its significant investment in advanced numerical simulation technologies.

1.1.2 Product Overview and Technical Approach

Put simply, V&Vis the process of ensuring a simulation is fit for its intended purpose. More precisely, Verification is the process of ensuring that a simulation is a faithful representation of the developers conceptual description and specifications (did I build the thing right?), and Validation is the process of determining the degree to which a model is an accurate representation of the physical model from the perspective of the intended use case (did I build the right thing?) ?. It is a long held belief of the software development community that the sooner a error is detected, the cheaper it is to fix. The same is true for numerical M&Sand as such, V&Vshould be an integrated component of the software development life-cycle. Some examples of the tasks involved in the V&Vof a numerical simulation include:

- Development of a V&Vplan
- Implementation of software development best practices (e.g. version control, unit and regression testing, code reviews, etc.)
- Mathematical and algorithmic testing (convergence analysis, mesh refinement studies, method of manufactured solutions, etc.)
- Development of a robust benchmark testing suite
- Uncertainty quantification and sensitivity analysis
- Comparison of simulation results with experimental data and results from third party simulations.
- Review of the implementation and results by experts in the field
- Documentation of the V&Veffort

In this proposal, we make the distinction between V&Va simulation package and the end-user V&Vof a solution obtained using a simulation package. These two processes are not independent; indeed, V&Vof a numerical simulation package almost always includes a set of verified and validated benchmark tests; likewise, the assertion that a package is verified and validated often forms the foundation of trust in the end user V&Vof a new simulation configuration. The focus of this proposal and the phase II project will be end-user validation; however, there is no reason that the functionality imparted by the proposed framework could not be used in the V&Vprocess of the overall simulation package. The key issues associated with end-user verification and validation, and the approaches that the VnV framework will take to address them are:

- **In-situ Testing And Analysis:** Unit tests are an extremely effective mechanism for ensuring a algorithm behaves as is expected. However, such testing is an unavoidably discrete process that, by definition, cannot cover every possible outcome. This fact is particularly true for numerical algorithms, where even small changes (e.g., input parameters, mesh geometry. etc.) can cause algorithms to diverge, or worse, converge to the incorrect solution. As such, a robust V&Vreport should not only include a description of unit tests completed, but also a detailed set of in-situ tests and assertions that run as the simulation proceeds. The VnV framework will include a sophisticated test injection system. At run time, the system will automatically detect, configure assimilate data obtained from injection points in any VnV equipt library that is linked to the executable. For example, once fully integrated at all levels of the simulation heirarchy, the user of a MOOSE tool would be able to detect, configure, modify and run tests at injection points defined in the source codes of hypre, PETSc, libMesh, MOOSE, etc. This cross library support will allow for in-depth, expert directed, end-user V&Vwithout ever needing to touch the source code of third party libraries. A key feature of the framework is that, while the injection points are hard coded into the source code, the tests can be added or removed from the simulation at runtime using a XML configuration file without the need to recompile. The tests themselves are compiled in external shared libraries loaded at runtime using the dynamic loader and could be as simple as asserting a value is positive, or as complicated as completing a statistical comparison between a distributed array and some experimental data.
- **Reusable Software Components:** While the specific details of V&Vvary from application to application, the macro scale algorithms used are relatively consistent, including; mesh refinement studies, using the method of manufactured solutions, sensitivity analysis, uncertainty quantification and error propagation. Many of these algorithms can, and in some cases already have (see, NIMROD, DAKOTA, MASA) be implemented as black-box or near black box solutions. The VnV framework will look to capitalize on this fact by providing a set of robust, near black-box, V&Vtools that can be integrated into codes using the VnV runtime test injection system. Where possible, these tools will leverage existing black box solutions such as the MASA repository for MMS analysis and the DAKOTA and NIMROD tools for UQ and SA.
- **V&VTesting Efficiency:** Performing a large number of V&Vtests in a distributed environment will be expensive, both computationally and due to the data movement required to deal with the domain decomposition employed by the application. Where possible, the VnV framework will offer functionality for offloading the execution of in-situ tests to external processes. Offloading of the tests to an external server will significantly reduce the run time costs of in-situ V&Vwith the VnV framework; further increasing its utility in HPC applications. The framework will also support job based parallelism for tests that require multiple simulations to be run (parameter optimization, mesh refinement).
- **Documentation Generation:** With software packages under almost constant development, and new and improved packages being released on a regular basis, keeping an up-to-date V&Vreport is in an almost impossible task. The VnV toolkit will include automatic VnV report generation in the form of a server-less html web page that can be viewed in any modern web browser. The entire web page will be customizable using an extended markdown format and will include support for standard

markdown formatting, latex formatting, images, videos, self-sorting tables, a two dimensional charts using charts.js and three dimensional visualization with paraview, vtk.js and three.js. The format of the overall web report will follow the templates outlined in the DoD TODO VV template files. These templates have been used over many years in the VV&Aof DoD simulations to great effect. Of course, additional custom templates will also be supported.

In summary, once integrated into an application, the VnV framework will provide a simple mechanism for creating self verifying, self describing, explainable numerical simulations. This will significantly reduce the burden associated with V&Vfor end user simulations, thereby increasing the usability of the tools for non-expert end-users. The framework will support a full range of in-situ V&Vtests with functionality for offloading in-situ tests to external processors. Finally, the package will include support for generating a highly customizable, web-based V&Vreport. A commercial license will be required to incorporate the VnV toolkit into for profit applications; however, core functionality of the VnV toolkit will be released as open source for use in academic and enterprise applications, with RNET providing commercial support, training and integration contracts to interested parties.

1.2 Anticipated Public Benefits

The initial customers of the VnV framework will include the businesses and other institutions (e.g., ANSYS, Cd-adapco, government labs, universities) that develop large-scale numerical modeling and simulations. To these customers, RNET will provide the training, support and integration services required to quickly and efficiently integrate the VnV framework into new and existing code bases, as well as contract based services for extending the toolkit to fit certain needs. For these customers, the benefits could include streamlining of the companies internal V&Vpractices, and an increase in the usability, reliability and confidence in their final product. However, the true beneficiaries of the VnV frameworks end-user V&Vfunctionality are the users of the advanced numerical simulation products. By removing the burden associated with V&V, the VnV framework will ensure erroneous errors do not propagate into final designs, while also opening the door to using simulations tools for which V&Vwas previous infeasible. All in all, the VnV framework will afford researchers and engineers with the knowledge required to drive the next generation of technological advancements.

1.3 Phase I and Feasibility Demonstration

The primary goal of the Phase I effort was to demonstrate the feasibility of the proposed approach to facilitate end-user verification and validation of numerical simulations. To that end, the project team spent the majority of the Phase I project developing an initial prototype that demonstrates the core functionality of the VnV toolkit. In particular, the team focused on implementing an initial version of the the injection point system and the automated report generation system.

1.3.1 VnV Injection Point System

The project team spent a large portion of the Phase I effort investigation approaches for implementing the run time injection point system, including automatic source code modification, custom llvm compilers and binary rewriting. In the end, the simplest approach, using hard-coded injection points, turned out to be the most portable, reliable and versatile approach. The injection point system developed in Phase I consists of four main components; the injection points, the tests, the input/output system and the runtime configuration module.

The core functionality of the injection point system was written in C++ with a focus on minimizing the amount of code required to integrate the injection points into existing code bases. As shown in Figure ??, integration of an injection point into an existing code is a simple, three step process; (1) include the “vv-runtime.h” header file, (2) place injection points in the code and (3) write the injection point specification.

At the core of the system is the INJECTION_POINT macro. The format for this macro is:

```
INJECTION_POINT( <unique-name>, <stage>, <type> <variable>, <type> <variable>, ...
```


During compilation, the the macro converts the injection point call to a function call of the form:

```
VV_injection(<unique-name>, <stage> , <count> , <type> , (void*) &<variable> , <type>
```

where count is an integer representing the number of variables and type is a string defining the type.

Here, the unique name represents the id that will be used to define the injection point in the configuration files and the final reports. This name must be unique across all injection points in an executable. The stage parameter is an integer value that defines the step that this injection point belongs to. Using this parameter, the developer can set up multi-staged injection point testing. Tests defined on staged injection points stay in scope across all stages allowing tests to collect data across multiple code locations. The variable parameters allows the developer to specify which variables can be inspected and tested at each injection point. The phase I prototype uses a combination of string based type checking, static casts to (void*) and C style variadic function calls to achieve this task. To enter a variable, the developer must pass the type and the variable to the macro. When the function is called at runtime, the VnV injection point factory compares the types accepted by the specified tests against the type specified at the injection point. There are several risks associated with this approach, the most important of which are avoiding memory corruption and ensuring constant correctness of the variables across tests. A key goal of the Phase II project will be to mitigate these risks and/or, to explore compiler based approaches to inspecting the desired parameters (see Section ??).

The final step in the creation of an injection point is to write the injection point specification. The specification is a YAML formatted file containing the content used to populate the final VnV report. Figure ?? shows a sample specification for the injection point shown in Figure ?. The only required parameter in the specification is the name; however, the more information that is entered here, the more informative the final report will be. The content section represents the text that will be shown in the final report each time the code reaches the associated injection point. The content variable is specified in an custom markdown format that supports all standard markdown commands, along with a range of data visualization extensions (described below). Content provided in the “sections” map will be listed as sub sections of the content parameter in the final report.

The second facet of the injection point system are the specific V&V tests. The development of the test interface was based on the idea that tests should be loaded at runtime and defined independently of the source code. To achieve this, the test interface was built using an interface-based C++ plugin pattern.

The first step in the development of a new VnV test is to create a testing library. The framework includes a python based test library generation script that will automatically build the directory structure and makefiles required to build this library. All that is required of the user for this step is to call this script with a unique name.

Once the library has been initialized, the user can begin to develop individual tests. Writing a test is as simple as implementing the IVVTest Interface shown in Figure ?. To assist in the development of tests, and to avoid issues with incorrect type-casting, a python based test generation script is created as part of the library initialization step. This script can be used to automatically generate the boiler plate code required to implement the testing interface while also taking care of the required typecasting. As shown in Figure ?, using this script significantly reduces the complexity of implementing a new test.

Implementing the IVVTest interface is a three step process. First, the developer must define the names and types of the parameters the test will support. This is completed in the constructor by adding the name and type of certain parameters to the parameters list. The only requirement on parameters is that the names should be unique within the test.

To ensure efficiency in HPC settings, the VnV framework uses ADIOS for all data output. Writing data inside a test using ADIOS is a two step process that requires the user to (1) declare the output variables that will be written and (2) write the data. Tests should declare the output variables in the DeclareIO function. Note that it is not a requirement that the output variables be defined in advance, however it is good practice and allows for more efficient handling of the meta-data in ADIOS.

The actual testing and writing of the data should occur in the runTest function. The variables passed to the test are direct pointers to the variables in the code. Hence, tests should not modify the pointers in any way; however, the phase I prototype does not strictly enforce this requirement. Beyond that requirement, there is not limit on the procedures that can be run inside a test.

As with the injection points, the final step in the test creation process is to define the test specification. This specification acts as the link between the data collected during testing and the final report. Figure ?? shows an example specification for a test that collects convergence factors from a PETSc linear solve routine. This specification uses the VnV extended markdown format (described below) to automatically generate a plot based on the convergence data extracted during the execution of the test.

Our vision for the VnV toolkit is that each numerical library will ship with hard-coded injection points and a custom VnV Testing library. The core VnV framework will act as a single interface to these individual libraries, allowing the end-users to build explainable numerical simulations with integrated, in-situ testing at every level of simulation hierarchy. Once the injection points have been declared and the tests created, the next step is to develop the XML test configuration file. Using this XML file, users can configure which testing libraries to load and which tests to run. Please see the final report for a full description of the XML format supported by the phase I prototype. With the configuration file in hand, performing VnV in a simulation is as simple as including a header file, linking the VnV library and calling the VnV Initialization function.

1.3.2 Automated Report Generation

In addition to the injection point system, the project team also developed an initial prototype for automatically generating the final report. After accessing the strengths and weaknesses of multiple different approaches, the project team decided on a server-less HTML/JS format generated using a custom extension of the markdown format. The primary benefit of this approach is portability - the report can be displayed in any web browser - but other benefits include interactive components, non-linear data presentation and high levels of customization. Moreover, the server-less nature of the HTML web-page allows for direct publication on any static web hosting service (github.io, AWS S3, etc.).

At the core of the report generation system is a custom markdown extension. This extension supports a range of custom data visualization features that interact directly with the output data in the ADIOS2 output file. For example, Figure ?? shows a markdown snippet for automatically generating an interactive two dimensional plot using time series data collected during VnV testing. Likewise, Figure ?? shows the markdown required to display a VTK .vtp file in the final report. The extension also includes support for custom post-processing scripts. This allows for infinite possibilities with regard to processing the test outputs. For example, it is possible to include a script that averages the array in a tests, and/or to run a paraview script that generates VTP files based on the test data.

Figures ?? show screen-shots of a sample VnV report generated using a set of toy testing libraries. The main layout consists of three components; the carousel, the index and the content. The carousel is an optional component that can be used to highlight important results of the simulation. It accepts up to ten pictures, each with its own custom caption. Simpler static headers are also possible. The index and content are generated automatically based on the VnV output file. Each node in the index represents an injection point encountered during the simulation. As such, this index represents a coarse grained view of the simulations call stack. At the top of each injection point section is the injection point content as specified in the injection point specification. Following this is the output of each test completed at the injection point. Last is the list of children. These children represent injection points that were initialized between the first and last step of a staged injection point. A live version of this sample VnV report will be available at <http://www.rnet-tech.net/VV/sample.html> until the date of award notification.

1.3.3 Demonstration in a Moose Application

To demonstrate the utility of the method, the project team placed several injection points in the main function of a MOOSE example “ex01”, one injection point in the PETSc Initialize function and one injection point in the PETSc Finalize functions. This is an extremely simple example that did not test the full limits of the new API; however, it did act to verify that the phase I prototype can be used to perform in-situ verification and validation in a across multiple libraries through a single interface. A screenshot of the final V&Vreport obtained from those tests is shown in the final report.

In summary, during Phase I, the project team created a functioning prototype of the VnV framework that provides:

- A clean mechanism for inserting injection points in existing codes
- A simple interface for defining custom tests
- A customizable approach for automatic post-processing of testing data and injection points
- A python based report generation code that automatically creates an interactive server-less web report based on the VnV output.

As will be described in the workplan, the goal of the Phase II project will be to take this initial prototype and extend, harden and optimize it such that it can be efficiently used in high performance computing applications.

2 The Phase II Project

2.1 Technical Objectives

The requirements being addressed include the development of a robust framework for in-situ verification and validation in general purpose numerical simulation packages. In particular, the objective of this project is to address the need for tools that automate verification of end-user numerical solutions in the NEAMS toolkit and workbench.

In Phase II, RNET Technologies and ORNL will pursue the following objectives:

1. {Harden and extend the core VnV functionality developed during phase I. In particular, the phase II effort will look to determine the optimal approach for implementing the run time configurable test injection system such that the risks associated with memory corruption and constant correctness are minimized. This objective will include the miscellaneous tasks required to prepare the framework for release, including integration with a unit testing framework, further development of the documentation generation engine, and documentation. }
2. {Develop mechanisms for efficient data movement in a distributed environment. This objective will look to determine and implement optimal approaches for comparing data stored in distributed arrays against an expected result stored on disk. The key issue here is to define approaches for describing the domain decomposition of the distributed array such that the experimental data can be distributed in an efficient manor. In-situ comparison of variables with experimental and/or analytical results will be a defining feature of the framework because it significantly reduces the amount of IO required in V&Vtesting, while also providing a fine grained mechanism for detecting at what point a solution diverges from the expected result. }
3. {Optimize test execution times. Initial development will focus on mechanisms for offloading tests to an external server. Offloading of tests to an external VnV testing service has the potential to significantly reduce overall runtime. This key issue to address here will be to develop a mechanism for offloading data such that the data transfer is faster than simply running the test in-situ. The initial focus will be on determining the best framework for offloading simple tests (MRNet, SNOBall, ADIOS streaming, etc.). After that has been implemented, the focus will shift toward implementing job based parallelism for tests the require the simulation to be run multiple times. }

4. { Develop a robust set of generic V&V tests. The development of these tests (e.g., mesh refinement studies, the method of manufactured solutions, sensitivity analysis, uncertainty quantification) will further equip the users with the tools required to robustly perform end-user V&V. The open research question the Phase II project will look to address is the optimal approach to integrating existing implementations of the tools (i.e., NIMROD, DAKOTA, MASA) into the in-situ V&V testing framework. Another open question is the development of a generic interface for mesh refinement studies such that we can automatically generate the required grid hierarchy. }
5. { Demonstrate the value of the VnV framework as a component in the NEAMS tools and into the NEAMS workbench. The true benefit of the VnV toolkit will only be realized if we can drive wide scale uptake across the entire numerical simulation community. By showing the toolkit can be used in the NEAMS tools, and in particular, MOOSE, libMesh and PETSc, we will demonstrate true potential of the product in libraries that are already considered cutting edge across the industry. Integration into the NEAMS tools will answer the NEAMS call for tools that support end-user verification of numerical simulations. Integration into the NEAMS workbench will provide access to the tools in the seamless manner users of the workbench have come to expect. }

2.2 Work Plan

As described in the objectives, the final deliverable of the Phase II project will be a fully functional, efficient, battle tested framework for integrating advanced end-user verification and validation into general purpose numerical simulation packages. In what follows, we outline the work required to satisfy the objectives outlined in the previous section and to achieve these goals.

2.2.1 Hardening and Optimization of the Injection Point System

The injection point system developed during Phase I uses C style Macros and string based (void*) pointer casts to declare the variables available for inspection at each injection point. When implemented correctly, this is an efficient (void* casts are almost free) and portable (it uses low level C functionality supported by all C/C++ compilers) approach. In Phase II, the project team will develop a custom compiler that supports an annotation based specification for the injection points and the injection point variables. This annotation based specification will be designed to address two key issues with the Phase I approach:

- **String based pointer casts:** The injection point system developed in Phase I requires the developer to provide a string that describes the type for each variable available at an injection point. Under the hood, the injection point system using string comparisons to ensure compatibility between test parameters and injection point variables. The key issue with this approach is that the strings specified at each injection point cannot be verified during compilation. This will cause major issues in cases where, say, the developer changes the type of a variable, but forgets to update the type string in the injection point. The custom source code processor developed in Phase II will automatically detect the type of the variables passed to the injection point system, removing the requirement for a hard-coded type string.
- **Restrictive type specifications:** The current system uses C compliant pre-processor macros to simplify the process of describing injection points and variables. The benefit of this approach is that the injection points can be compiled into any application without significant changes to the build system. The downside is that single-pass, text based macros place a significant restriction on the functionality that can be implemented. The annotation system developed in Phase II will be far more dynamic, allowing users to have full control over what variables can be accessed at each injection point, including the ability to provide access to internal components of data structures, describe the domain decomposition of distributed arrays and to complete pre-test processing of variables. The annotation system will also provide a mechanism for suggesting default tests to be run at each injection point, and will provide a mechanism for injection point detection in non-object oriented programming languages where runtime injection point detection cannot be completed.

This compiler will be written using Clang. Clang is a compiler front end for the C family of languages. It is a well supported, well documented compiler package designed as a drop in replacement for GCC. In particular, Clangs simple API for defining custom Pragma routines, combined with its robust API for walking the Abstract syntax tree of the code make it the perfect choice for developing this annotation based system for defining injection points. Clang has seen wide scale uptake across the software development world and, in recent years, has emerged as a realistic competitor to the ubiquitous GCC compiler suite.

Figure ?? shows how the annotation based injection point system will look in a C/C++ code. The user will annotate simple variables with the “@Variable(“IP_1“, “IP_2“, ...)” annotation. In that annotation, users will be able to specify the injection points for which that variable is available. Additionally, users will be able to specify additional parameters that define how the variable should be accessed, e.g., for providing access to a member of a struct (line TODO), or for providing access to a certain element in an array (line TODO). For distributed arrays, the annotation based system will allow users to define approaches for obtaining the global and local ownership for each processor. Together, this annotation based system will allow users to fully control the access to these variables. The original Phase I approach will still be available for users that cannot switch to the new Clang compiler. Where possible, the custom compiler will make use of available C++11 RTTI information (dynamic_cast, typeid, etc.); however, a general C approach will be favored due to the large number of C programs still regularly used in scientific computing.

2.2.2 Efficient Statistical Comparisons in Distributed Settings

One of the major benefits of the VnV framework is that it allows for in-situ testing and analysis in distributed systems. This direct, in-situ access allows for a detailed level of testing that would otherwise not be feasible in V&V approaches that require writing data to disk. For example, this in-situ analysis could allow the user to directly assert that all the elements in an array are positive, or that all the elements in an array are the same as the elements in another array without ever writing the data to disk.

The goal of the Phase II effort will be to the main issues associated with performing such tests in a large-scale distributed environment; determining the data decomposition and efficiently comparing arrays in a distributed setting.

In order for the testing algorithms to act of the data structures, they must be aware of the way that the code has distributed the code across the parallel processors. As such, a key goal of the Phase II effort will be the implement efficient, user friendly approaches for describing the domain decomposition for arrays with both regular and irregular decompositions. For regular decompositions, the approach taken will be to allow the user to specify a BlockMap function as described in .Blockmap follows the HPF syntax [21], and includes a collection of expressions that describe the decomposition scheme. In the blockMap function, each dimension of the data is classified as either “Block”, “Clyclic” or “Ignored”; where block decomposition means that a contiguous number of elements in the given dimension are mapped to different processes; Cyclic decomposition means each element in the given dimension is mapped to a different process in order; and ignored means that decomposition is ignored along that dimension. Figures ?? show some examples of typical decompositions for 2D arrays supported by the blockMap function. To support irregular decompositions, the Phase II effort will implement a generic algorithm for forming the global partition. Under the assumption that every processor knows its own local ownership range (in global co-ordinates), the global partition can be formed using a simple global collective communication (All-To-All communication). Once each processor knows the global partition, figuring out which processors own the required indecies is a simple task. There are some concerns with this approach, namely that it requires $O(P)$ storage and $O(\log(P))$ communication. Such costs are acceptable at modest processor counts, but could potentially become an issue in large scale (100K+ processors) settings. If the need arises, the project team will investigate implementing more advanced methods for determining intra processor communication patterns in a distributed data environment, including creating a distributed directory ? or through an assumed partition ?.

Once the global partition has been obtained, the important task becomes determining TODO.

2.2.3 Reducing Run-times with Test Offloading and Job Parallelism

Performing a large number of V&V tests in a distributed environment will be expensive, both computationally and due to the data movement required to deal with the domain decomposition employed by the application. A key goal of the Phase II effort will be to investigate and implement approaches for distributing tests for offloading tests for execution on separate processes.

The issue of efficient data transfer will be addressed using the ADIOS 2 Sustainable Staging Transport (SST). The SST is a classic streaming data architecture, that allows for direct connection between data producers and data consumers via the ADIOS2 write/read API. In HPC environments, SST uses the RDMA interconnects to ensure fast transfer of data between HPC applications; however, socket based connections are also supported. Due to issues associated with serialization of generic data-structures, the Phase II implementation will restrict data offloading to tests that work with the basic data types supported by ADIOS2 (strings, floats, double, arrays, etc). Generic data structures will be investigated in Phase III. The development of test offloading will proceed in two stages;

- The first step will be to develop the interfaces and annotations required to transfer data to external processes for testing. Test offloading will only be effective in situations where the cost of completing the required tests is large in comparison to the time required for transferring the data. To address this issue, the project team will implement a heuristic algorithms that determines the appropriate action based on the size and type of the data structures to be transferred and the type and number of tests to be executed. More robust approaches will be investigated after testing and analysis of the initial implementation.
- The second step will be to develop the separate VV executable that consumes the data and coordinates the execution of the required tests. As shown in Figure ??, this will be a MPI application consisting of one master processor and any number of slave processors. The role of the master processor is to efficiently allocate the incoming tests and data to the slave processors. This will include determining the appropriate number of processors to allocate to each test given the available resources, forming the required communication groups, and executing the tests. The role of the slave nodes will be to execute the tests as directed.

In many cases, it will be more cost efficient to execute the tests as separate jobs. To support this use case, the external test executable will support submitting tests for execution on local and remote machines using the Eclipse Parallel tools platform (PTP). PTP was developed by our collaborator as a simple unified interface for interacting with various job schedulers available in the HPC community (PBS, SLURM, etc.). PTP has seen wide-scale usage across the scientific community, and is used in such cases as TODO.

2.2.4 Integrate Third Party Tools for Mesh Refinement, UQ and Sensitivity Analysis.

Mesh refinement studies, Uncertainty quantification and sensitivity analysis are all essential components of a robust V&V regimen. To that end, the Phase II effort will include the development of V&V tests that integrate third party tools to complete these tests.

The UQ and SA tests will be developed using DAKOTA. DAKOTA provides a set of black-box tools for performing parameter optimization, UQ and SA. DAKOTA provides two interfaces for integrating these tools into applications; a black-box approach that uses preprocessors, postprocessors and the file system to complete the tests; and library API for direct, hard-coded testing. The Phase II effort will focus on integrating the DAKOTA tools through the library API as it allows for the tools to be applied directly to any independent function. To that end, the Phase II effort will include the development of an custom VnV testing library for direct integration with the DAKOTA library API. This will include the development of a flexible interface for setting up and running the DAKOTA tests, and the development of the test specification files for displaying the test results in an informative and interactive way.

Support for Mesh refinement will be developed as part a separate plug-able VnV testing library. The intended functionality of this testing suite will be to enable the completion of both uniform and adaptive mesh refinement studies given only the initial coarse mesh. This will require the development of methods for marking the mesh for refinement and methods for refining the mesh.

2.2.5 Integration with NEAMS tools and the NEAMS workbench

The final step in the Phase II project will be to integrate the VnV framework into the NEAMS tools. Doing so will allow us to demonstrate the performance of the toolkit in real codes with real applications; while also answering the solicitation call for tools that support end-user verification of numerical solutions in the NEAMS tools.

Integration into NEAMS tools will be a three step process:

- The first step will be to insert injection points into key locations in the MOOSE, libMesh and PETSc source codes. The optimal locations for inserting these injection points will be determined after detailed profiling of example codes; however, some obvious options include during each linear solver iteration, during each nonlinear iteration, during finite element matrix construction (if it exists) and inside the function for calculating the action of the Jacobian on a vector.
- The second step will be to allow users to configure the VnV testing directly in the MOOSE input file, using the MOOSE input file syntax. This will provide users of MOOSE with a seamless mechanism for setting up and running tests. This will be completed by developing a custom MOOSE “Action” that processes the input at runtime to setup the core VnV testing functionality.
- The third step will be to enable context-aware auto-complete for MOOSE based VnV configuration files in the NEAMS workbench. The workbench has integrated support for extracting the information required to setup input file verification and auto-complete from MOOSE applications; however, there will likely be some additional work required to correctly setup the auto-complete and verification for tests stored in external libraries. In particular, the current system requires that the tests adhere to a specific XSD specification for configuration, but there is not yet a system for extracting the exact parameters required to inject a specific test. To do this, we will create a lightweight executable that loads external test libraries and extracts the appropriate parameter specifications in the “SON” format supported by workbench.
- The last step will be the development of an interface for editing the YAML specification files used to generate the final report and an interface for viewing said final reports. The GUI will be written using standard HTML and Javascript and then integrated into the workbench using the QT QWebEngineView component (QWebView is the workbench is using QT 4). The QWebEngineView component allows for interactions with the core QT window (i.e., open, save, copy, paste, etc.) however, it is important to note that this approach will not allow for a truly “native” experience within the workbench. The decision to develop in JS and HTML rather than native QT was made to ensure the final product, the VnV framework, is consistent, and can be used in applications outside of the NEAMS lines of tools. Every effort will be made to ensure the GUI conforms to the NEAMS workbench standards and specifications.

2.3 Performance Schedule and Task Plan

RNET would like to present the project ideas and research plan to the DOE Program Manager and other interested scientists. The meeting will be used to discuss features, and identify the specific NEAMS applications and computer resources that will benefit from this project. This meeting will be scheduled soon after the Phase II contract is awarded. The meeting can be hosted at RNET, a DOE site suggested by the Program Manager or via a teleconference.

Figure 1: Overview of task dependencies and timeline.

RNET will submit all reports as required by the contract (e.g., annual reports, a continuation report, summary reports, and a final report) to the DOE program manager and other interested DOE scientists.

The research and development topics described in Section ?? will be addressed by the tasks described in the remainder of this section. Most tasks require active collaboration between RNET and its collaborators. Figure ?? summarizes at a high level the dependencies among tasks and approximate anticipated task durations. The duration of the Phase II project is 104 weeks. Specific details are included in the description of each task.

2.3.1 Task ??: Hardening of the Phase I VnV Injection Point System.

In this task, the project team will develop an efficient, intuitive and type-safe injection point system for specifying testing points in, and across, advanced numerical software packages. Building on top of the Phase I prototype, this work will look to implement the injection point system in a way that allows for in-situ variable inspection in a robust and consistent way. The two main that need to be addressed are (1) the distinct lack of a robust reflection API in C++ makes defining functions that support generic arguments of unknown classes in a type safe way difficult, (2) the overhead associated with V&V injection points in cases where no testing is requested and (3) the difficulties in ensuring tests do not modify the variables in any way. Implementing this will require looking into the possibilities of implementing a LLVM compiler option for run-time removal of unused injection points, and/or the development of binary rewriting tools and the integration of parallel relative debugging techniques for detecting un-authorized and/or unintended changes is the run-time variables.

{RNET will work on the implementation for this task and ORNL will provide inputs and guidance.}

2.3.2 Task ??: Develop methods for offloading tests to external processes.

Performing a large number of V&V tests in a distributed environment will be expensive, both computationally and due to the data movement required to deal with the domain decomposition employed by the application. In this task, the project team will investigate the development of a mechanism for the evaluation of V&V tests that is lightweight and minimizes resource consumption. The initial approach to this will be to determine the requirements for efficient V&V test execution, and assess the capabilities provided by frameworks such as MRNet, SNOflake, ADIOS, etc. to determine if they can meet these requirements. Depending on the results of this assessment, work will be undertaken to create a library that can be integrated with the V&V framework. This library will be based on either one of these frameworks or using a custom solution. The project team will also determine the feasibility of integrating the Parallel Tools Platform (PTP) as a mechanism for minimizing V&V runtime through job parallelism. If appropriate, the team will extending PTP to fit this purpose.

{RNET will work on the implementation for this task and ORNL will provide inputs and guidance.}

2.3.3 Task ??: Development of efficient statistical V&V tools with a focus of performance in large scale distributed settings.

In this task, the team will perform R&D work to determine and implement efficient statistical based V&V tools that can be applied in distributed settings. Consider that case of comparing the simulated solution to some known experimental results. In a distributed setting, the user must either distribute the experimental data across the nodes based on some known description of the data decomposition (difficult in a general setting), or gather to solution to some root node (expensive, and in many cases impossible. In this task, the project team will look into efficient mechanisms for inferring, detecting, and or describing the data decomposition such that this style of tests can be completed in an efficient manner.

{RNET will be responsible for this task. ORNL will provide guidance on developing the framework for ORNL CADES.}

2.3.4 Task ??: Development of Generic tools for Mesh refinement, Uncertainty quantification and Sensitivity Analysis.

In this task, RNET will implement generic V&V tools for mesh refinement, uncertainty quantification and sensitivity analysis. In the case of mesh refinement, the approach taken will be to create a generic interface for interacting with the automatic mesh refinement functionality that already exists in finite element libraries like LibMesh and Fenics. The overall goal is to create a generic VnV test that can be attached to the main function of the executable such that it automates the process of running mesh refinement and mesh convergence studies. In the case of UQ and SA, the project team will develop an interface for specifying the tools available in DAKOTA as VnV tests.

{RNET will be responsible for this task. ORNL will provide guidance on developing the framework for ORNL CADES.}

2.3.5 Task ??: Extension of the VnV report generation system

In this task, the project team will complete the development of the VnV report generation system. The primary goal of this task will be to provide support for generating VnV reports that conform to the specification outlined in the DoD VV report templates. The task will also include a full implementation of the VnV markdown extension to support a wide variety of data visualization components, the development of the interfaces required for displaying unit and regression testing reports, and the software indirections required for assimilating multiple VnV reports into a single document that can be used in the V&V of an entire simulation package.

{RNET will work on the implementation for this task}

2.3.6 Task ??: Integration into real applications, including the NEAMS Tools

{ In this task, the program team will integrate the VnV framework into a variety of real applications. Initially, this testing will be completed in tools used heavily in the NEAMS toolkit; MOOSE, libMesh and PETSc, but other third party applications will also be investigated. To ensure seamless integration with the MOOSE tools, the project team will reimplement the XML based configuration file using the MOOSE input file format. This will allow for the configuration of the VnV tests in a MOOSE application directly from the input file.

The goal of this task will be to generate informative, production quality VnV reports for a number of examples available in the MOOSE testing suite. Doing so allows us to test every facet of the proposed framework, while also acting as the first demonstration of the value provided by the framework. These results of these tests will be hosted on the RNET website as they become available.

}

{ RNET will be responsible for this task and ORNL will provide guidance on various technical implementations and details. }

2.3.7 Task ??: Development of an interface for the NEAMS workbench

{ In this task, the project team will integrate the toolkit directly into the NEAMS workbench. This will be a two stage process. First, the project team will implement the required interface files for enabling the context aware auto-complete features available in the NEAMS workbench for the MOOSE based configuration file specification developed in the previous section. MOOSE based input files are already largely supported in the workbench; however, there will likely be some issues with determining which tests are applicable at which injection points. Second will be developing an interface for customizing and viewing the final VV report. As part of the Phase I effort, the project team demonstrated viewing the final V&V report in a QT WebView component. The workbench is also built on top of QT, hence we do not expect too many difficulties on that front. Instead the key objective will be to develop the mechanisms for displaying customizations made to the final report in real-time within the NEAMS workbench. }

{ RNET will be responsible for this task }

2.4 Facilities/Equipment

2.4.1 RNET Facilities

RNET has the necessary office equipment to manage an SBIR/STTR contract including networks, workstations, and accounting software. In addition, RNET has the tools (software and hardware) to evaluate and develop the technologies proposed here.

RNET currently has 9 development computers and a 10-node development cluster that can be used for development and testing in this effort. Each cluster node has two quad-core or hexa-core XEON CPUs, 24-32GB of DRAM, 500+GB of local disk. Two data networks are available, a COTS 1 Gbps Ethernet network and a 10 Gbps Ethernet network. The Linux development nodes and the RNET cluster have the necessary Linux/GNU toolchains and development environments including; GNU tool chain, Microsoft .Net Framework, and Java Standard Edition.

2.4.2 ORNL Facilities

The Oak Ridge National Laboratory (ORNL) hosts three petascale computing facilities: the Oak Ridge Leadership Computing Facility (OLCF), managed for DOE; the National Institute for Computational Sciences (NICS) computing facility operated for the National Science Foundation (NSF); and the National Climate-Computing Research Center (NCRC), formed as collaboration between ORNL and the National Oceanographic and Atmospheric Administration (NOAA) to explore a variety of research topics in climate sciences. Each of these facilities has a professional, experienced operational and engineering staff comprising groups in high-performance computing (HPC) operations, technology integration, user services, scientific computing, and application performance tools.

ORNL also has the Compute and Data Environment for Science (CADES) which is a fully integrated infrastructure offering compute and data services for researchers lab-wide. We will work with appropriate program managers to apply for allocation requests as appropriate.

The ORNL computer facility staff provides continuous operation of the centers and immediate problem resolution. On evenings and weekends, operators provide first-line problem resolution for users with additional user support and system administrators on-call for more difficult problems. ORNL also has state-of-the-art visualization facilities that can be used on site or accessed remotely.

3 Consultants and Subcontractors

Oak Ridge National Laboratory(ORNL) is the Research Institution for this proposal and will serve as a subcontractor for this SBIR.

Gregory Watson, PhD (Monsash,2000), is a Senior Research Scientist in the Computer Science Research Group at Oak Ridge National Laboratory. Dr. Watsons research interests include programming tools and development environments for high performance and scientific computing, software engineering practices, reproducibility, and education and training for scientists. In particular, Dr. Watson is the founder of the Eclipse Parallel Tools Platform, a project that was originally started as a collaboration between Los Alamos National Laboratory and IBM in 2004, and that continues to be used across laboratories, academia, and industry. He is also a founding member of the Eclipse Science Working Group, and project leader of the Eclipse Science Top Level Project. In particular, Dr. Watson has considerable experience developing and implementing efficient parallel debugging software for high performance computing environments and has wealth of experience in the development of highly scalable tools and communication frameworks for petascale high performance computing systems.

4 Principal Investigator and other Key Personnel

4.0.1 Ben O'Neill, PI

Ben O'Neill is a Research Scientist at RNET and will be the PI on this project. Ben is a full time employee at RNET and has sufficient time to dedicate to this project. Ben is a permanent resident in the USA (citizen

of New Zealand). The proposed work does not include any Export Control or ITAR restriction, as such, this work visa should be sufficient. In addition to the current project, Ben is the lead developer in RNET's Cloudbench project which aims to develop a web-enabled interface for remote execution and visualization for nuclear physics tools. Ben was also heavily involved in the development of RNETS SolverSelector framework for facilitating automatic linear solver selection in high performance applications through machine learning. His background is in Applied mathematics with a focus on high performance computing and parallel-time integration. His thesis work involved the optimization and implementation of a parallel time integration codes for nonlinear PDEs including implementing a fully adaptive and parallel space-time solver using the Fenics finite element package and implementation of $O(\log(n))$ temporal load balancing scheme.

4.1 Gerald Sabin

Dr. Gerald Sabin, Project Manager at RNET, will be the senior advisor for this project. Dr. Sabin is a full time employee of RNET, and has sufficient time to dedicate to project tasks as indicated in the cost proposal. Since he is a US Citizen, he can undertake relevant integration work in Export-Controlled areas of the project, if necessary. Currently, he is working on several Scientific Computing (HPC) SBIR/STTR projects at RNET. He is the PI for this Phase II SBIR Cloudbench project (DE-SC0015748) and the ongoing Phase II DOE SBIR for the Automated Solver Selection for Nuclear Engineering Simulations. He has also worked on distributed memory, GPU, multi-core and SIMD optimizations to the Air Force's Kestrel code (DOD Contract#:FA9550-12-C-0028) and has also been involved in developing fine-grained power profiling hardware and software tools for HPC application profiling (DOE Contract#:DE-SC0004510). He has also been the PI on several other related projects including a NASA Phase I project developing SIMD optimizations for Monte Carlo codes (NNX14CA44P), developing parallelization optimizations for PETSc (DOE Contract#: DE-SC0002434), developing data virtualization support and bitmap indexing for massive Climate Modeling data sets (DOE Contract #:DE-SC0009520), and developing the SmartNIC software stack for application-aware network offloading (DOE Contract#: DE-FG02-08ER86360).

5 Related Work

RNET and ORNL have past and current experience in several SBIR/STTR projects on modeling and simulation, high performance computing, and large data formats. Some of these projects are briefly described below.

5.1 RNET's Related Work

5.1.1 Automated Solver Selection for NEAMS Tools

RNET in collaboration with University of Oregon (Dr. Boyana Norris) is developing an add-in feature for the NEAMS toolkit being developed by the Department of Energy. This work is being done as part of DOE Phase II STTR project (Contract Number DE-SC0013869). This add-in feature being developed by RNET leverages machine learning techniques to automatically select the optimal solver based on run-time dependent features of the problem and the underlying compute architecture with minimal runtime overhead in solver selection during the course of NEAMS simulations.

5.1.2 Cloudbench

5.1.3 Scaling the PETSc Numerical Library to Petascale Architectures

RNET has developed an extended version of the numerical library PETSc ? in collaboration with Ohio State University and Argonne National Lab. PETSc is an MPI-based numerical library of linear and nonlinear solvers that is widely used in a variety of scientific domains. With the emergence of multicore processors and heterogeneous accelerators as the building blocks of parallel systems, it is essential to restructure the PETSc code to effectively exploit multi-level parallelism. Changes to the underlying PETSc data structures are required to leverage the multicore nodes and GPGPUs being added to the "cluster architectures".

This project was funded by Department of Energy under the STTR program from August 2010 (Contract Number DE-SC0002434) to May 2013. Dr. P. Sadayappan (OSU) and Dr. Boyana Norris (ANL) have played a key role in this effort by serving as technical advisors. As part of the project, the team has investigated ways for the PETSc library to fully utilize the computing power of future Petascale computers. Novel sparse matrix types, vector types, and preconditioning techniques that are conducive for GPU processing and SIMD parallelization have been integrated into the PETSc library. The matrix vector operations have been optimized for specific architectures and GPUs by utilizing the autotuning tools.

5.1.4 Optimization of Kestrel for Emerging Architectures

RNET and OSU are performing this work as part of an active DOD Phase II STTR (FA9550-12-C-0028, Highly-Scalable Computational-Based Engineering). Based on the identification of the main performance bottlenecks in the Kestrel CFD code (based on the AVUS CFD solver), we are developing enhancements to improve the performance of the kCFD solver, as well as interface other scalable Krylov subspace sparse solver libraries to Kestrel. The proposed work will address the effective exploitation of parallelism at multiple levels: SIMD/SIMT level, multi-core level, and multi-node level.

As part of this project CUDA kernels are also being explored for the bottlenecks in the CFD and CSD solvers. For instance, a GPU-based CFD solver with an identical interface to the current Block-Seidel solver is being explored.

5.2 ORNL's Related Work

TODO

A solid foundation has been laid to achieve the objectives of the proposed project. The final piece of the puzzle is the optimization and hardening of these core approaches for efficient execution in real applications. Once this is accomplished, it will be possible to seamlessly integrate functionality for end-user V&V into any general purpose numerical simulation software. The collaboration between ORNL and RNET is a perfect fit to facilitate this final step. RNET brings to the table the required computer science expertise to satisfy the needs of this project, as is evident from a description of their related work, and ORNL brings a wealth of experience developing novel HPC software solutions, developing robust debugging tools for use in HPC environments and working with a range of real application codes.